

In [1]:

```

1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
7 from sklearn.linear_model import Ridge, RidgeCV, Lasso
8 from sklearn.preprocessing import StandardScaler

```

In [2]:

```

1 df=pd.read_csv(r"C:\Users\HP\Downloads\Advertising.csv")
2 df

```

Out[2]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [3]:

```
1 df.head()
```

Out[3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [4]:

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio        200 non-null    float64
2    Newspaper    200 non-null    float64
3    Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [5]:

1 df.describe()

Out[5]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

In [6]:

1 df.tail()

Out[6]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [7]:

1 df.shape

Out[7]:

(200, 4)

In [8]:

```
1 df.isna().any()
```

Out[8]:

TV False  
Radio False  
Newspaper False  
Sales False  
dtype: bool

In [9]:

```
1  
2  
3 plt.figure(figsize = (10, 10))  
4 sns.heatmap(df.corr(), annot = True)  
5
```

Out[9]:

<Axes: >




In [10]:

```

1 df.drop(columns = ["Radio", "Newspaper"], inplace = True)
2 #pairplot
3 sns.pairplot(df)
4 df.Sales = np.log(df.Sales)
5

```



In [11]:

```

1 features = df.columns[0:2]
2 target = df.columns[-1]
3 #X and y values
4 X = df[features].values
5 y = df[target].values
6 #split
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8 print("The dimension of X_train is {}".format(X_train.shape))
9 print("The dimension of X_test is {}".format(X_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)

```

The dimension of X\_train is (140, 2)

The dimension of X\_test is (60, 2)

In [12]:

```
1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(X_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(X_train, y_train)
10 test_score_lr = lr.score(X_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0  
The test score for lr model is 1.0

In [13]:

```
1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(X_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(X_train, y_train)
10 test_score_lr = lr.score(X_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0  
The test score for lr model is 1.0

In [14]:

```
1 #Ridge Regression Model
2 ridgeReg = Ridge(alpha=10)
3 ridgeReg.fit(X_train,y_train)
4 #train and test scorefor ridge regression
5 train_score_ridge = ridgeReg.score(X_train, y_train)
6 test_score_ridge = ridgeReg.score(X_test, y_test)
7 print("\nRidge Model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
10
```


Ridge Model:

The train score for ridge model is 0.990287139194161

The test score for ridge model is 0.9844266285141221

In [15]:

```
1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
3 #plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
4 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
5 plt.xticks(rotation = 90)
6 plt.legend()
7 plt.show()
```



In [16]:

```
1 #Lasso regression model
2 print("\nLasso Model: \n")
3 lasso = Lasso(alpha = 10)
4 lasso.fit(X_train,y_train)
5 train_score_ls =lasso.score(X_train,y_train)
6 test_score_ls =lasso.score(X_test,y_test)
7 print("The train score for ls model is {}".format(train_score_ls))
8 print("The test score for ls model is {}".format(test_score_ls))
9
```

Lasso Model:

The train score for ls model is 0.0


The test score for ls model is -0.0042092253233847465

In [17]:

```
1 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[17]:

&lt;Axes: &gt;





In [18]:


```
1 #Using the linear CV model
2 from sklearn.linear_model import LassoCV
3 #Lasso Cross validation
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
5 #score
6 print(lasso_cv.score(X_train, y_train))
7 print(lasso_cv.score(X_test, y_test))
8
```

0.9999999343798134

0.9999999152638072

In [19]:

```
1 #plot size
2 plt.figure(figsize = (10, 10))
3 #add plot for ridge regression
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
5 #add plot for lasso regression
6 #plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
7 #add plot for linear model
8 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
9 #rotate axis
10 plt.xticks(rotation = 90)
11 plt.legend()
12 plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
13 plt.show()
14
```



In [20]:

```

1 #Using the linear CV model
2 from sklearn.linear_model import RidgeCV
3 #Ridge Cross validation
4 ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
5 #score
6 print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train))
7 print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))

```

The train score for ridge model is 0.999999999997627

The train score for ridge model is 0.9999999999962467

In [21]:

```

1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(X,y)
4 print(regr.coef_)
5 print(regr.intercept_)

```

```

[0.00417976 0.          ]
2.026383919311004

```

In [22]:

```

1 y_pred_elastic=regr.predict(X_train)
2

```

In [26]:

```

1 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
2 print("Mean Squared Error on test set",mean_squared_error)

```

Mean Squared Error on test set 0.5538818050142158

Type *Markdown* and LaTeX:  $\alpha^2$

In [ ]:

1

In [ ]:

1

In [ ]:

1

In [ ]:

1

In [ ]:

1	
---	--

In [ ]:

1	
---	--

In [ ]:

1	
---	--

In [ ]:

1	
---	--

In [ ]:

1	
---	--